Architectural Integration of Genomic Multi-Omics into the SarcomaAI Framework: A Multimodal Deep Learning Approach for Precision Orthopaedic Oncology

The clinical management of soft tissue sarcomas (STS) remains a profound challenge in oncology due to a combination of rarity, extreme histological diversity, and highly variable biological behavior. Representing less than 1% of all adult malignancies yet comprising more than 70 distinct histological subtypes, sarcomas present a diagnostic and prognostic landscape that defies standard "one-size-fits-all" treatment paradigms.[1] Historically, clinicians have relied on clinical variables—such as tumor size, depth, and histological grade—to predict patient outcomes. While tools like the Sarculator nomogram and PERSARC have provided a degree of evidence-based guidance, they are fundamentally limited by their reliance on aggregate median data, which inevitably fails to capture the granular, individual variability inherent in mesenchymal tumors.[1] The development of the SarcomaAI framework marks a significant evolution in this field by employing a multimodal neural network (MMNN) that integrates clinical variables with three-dimensional magnetic resonance imaging (MRI) data.[1] However, even with the high-resolution morphological data provided by pre-treatment T1 post-contrast and T2 fat-saturated MRI sequences, a substantial portion of the prognostic variance remains unexplained.[1] The integration of genomic data—including transcriptomic profiles, copy number variations (CNV), and somatic mutations—offers the most promising pathway toward bridging this gap and enabling the subtype-specific precision that is required to improve the plateaued 65-70% five-year survival rates.[1]

The Foundational SarcomaAI Architecture and Gradient Blending

To supplement the SarcomaAI tool with genomic data, one must first analyze the technical constraints and capabilities of the existing end-to-end multimodal neural network. The current architecture is designed to handle two distinct data streams: a clinical subnetwork and an image subnetwork.[1] The image subnetwork utilizes a two-channel DenseNet-121 architecture, which is specifically optimized for 3D volumetric data. This network accepts MRI slices resized to $128 \times 128 \times 128$ voxels, employing a series of dense blocks where each layer is connected to every other layer within the block, facilitating feature reuse and reducing the risk of vanishing gradients.[1] This morphological processing identifies salient features such as tumor heterogeneity, vascularity, and necrotic cores, which are often indicative of higher-grade, more aggressive disease.[1] Parallel to this, the clinical subnetwork processes 11 key variables—age, sex, tumor location, diagnosis, neoadjuvant chemotherapy, tumor size, tumor volume, tumor depth, grade, metastatic status at presentation, and radiotherapy type—through a multilayer perceptron (MLP) featuring linear activations, batch normalization, and dropout to prevent overfitting.[1]

The integration of these streams is managed through a late-fusion concatenation mechanism followed by a final fully connected layer that outputs predictions for survival and metastatic risk.[1] A critical innovation within this framework is the application of gradient blending during the training phase. Standard joint training of multimodal models often suffers from "modality laziness," where the network converges prematurely on the easiest-to-learn modality—typically the low-dimensional clinical variables—at the expense of optimizing the more complex, high-dimensional data like 3D images.[1] Gradient blending solves this by monitoring the overfitting-to-generalization ratio (OGR) of each branch. The loss function is dynamically weighted according to the formula $L = w_{MM} L_{MM} + w_{img} L_{img} + w_{clin} L_{clin}$, where the weights w_i are adjusted periodically based on the validation set performance.[1] This ensures that both subnetworks are optimally trained, a principle that must be strictly maintained when introducing a third, even higher-dimensional genomic modality.

Feature Category	Variable Detail	Data Structure
Clinical Input	11 variables (Age, Size, Grade, etc.)	1D Tabular Vector [1]
Imaging Input	3D MRI (T1-weighted and T2 fat-sat)	$2 \times 128 \times 128 \times$ Voxel Grid [1]
Image Backbone	DenseNet-121 (Self-supervised pre-training)	3D Convolutional Neural Network [1]
Clinical Backbone	Multi-Layer Perceptron (MLP)	Dense layers with ReLU and Dropout [1]
Fusion Layer	Late Blending	Concatenation of latent vectors [1]

The Genomic Imperative: Molecular Drivers of Sarcoma Progression

The rationale for supplementing the SarcomaAI tool with genomic data lies in the fundamental biological differences between sarcoma subtypes. Unlike most epithelial carcinomas, which are characterized by high rates of point mutations, sarcomas are broadly divided into two genomic classes: simple-karyotype sarcomas and complex-karyotype sarcomas.[3] Simple-karyotype sarcomas, such as synovial sarcoma or Ewing sarcoma, are often driven by a single, pathognomonic chromosomal translocation or fusion gene. In contrast, complex-karyotype sarcomas, including undifferentiated pleomorphic sarcoma (UPS), leiomyosarcoma (LMS), and myxofibrosarcoma (MFS), exhibit massive genomic instability, widespread copy number alterations (CNAs), and low mutational loads.[3]

The TCGA-SARC project, which analyzed 206 adult soft tissue sarcomas across six major types, identified that somatic mutations are recurrent in only a handful of genes: $TP53$, $ATRX$, and $RB1$.[3] However, the prognostic signal found in copy number changes is profound. For instance, dedifferentiated liposarcomas (DDLPS) are characterized by significant amplifications in the 12q13-15 region, affecting genes like $CDK4$ and $MDM2$, which directly correlate with clinical outcome and therapy resistance.[3] Furthermore, transcriptomic signatures like the 67-gene Complexity INdex in SARComas (CINSARC) have proven to be more accurate predictors of metastatic relapse than the traditional French Federation of Cancer Centers (FNCLCC) histological grade.[11] These molecular markers provide a biological "ground truth" that morphology alone cannot fully reveal. Integrating these genomic signatures into the SarcomaAI architecture allows the model to reconcile visual indicators of aggression, such as peritumoral edema, with underlying genetic drivers of cellular proliferation and chromosomal instability.[11]

Genomic Data Type	Relevancy to SarcomaAI	Key Markers in STS
Transcriptomics	Reflects current cellular state and pathway activity.	CINSARC (67-gene signature) [11]
Copy Number (CNV)	Primary driver in complex-karyotype subtypes.	CDK4, MDM2, TP53, RB1 amplifications/deletions [3]
Somatic Mutations	Identifies specific oncogenic drivers.	TP53, ATRX, RB1 point mutations [3]

Epigenetics	DNA methylation reveals immune microenvironment.	T-cell and B-cell infiltration scores [4]
miRNA Profiling	Associated with survival in LMS and DDLPS.	Subtype-specific regulatory networks [4]

Designing the Genomic Subnetwork for High-Dimensional Data

The primary technical challenge in supplementing the SarcomaAI architecture is the integration of high-dimensional genomic data—often exceeding 20,000 transcriptomic features—into a model trained on a relatively small cohort of patients.[1] To prevent the model from overfitting to noise in the genomic data, the proposed genomic subnetwork must employ advanced dimensionality reduction and feature extraction techniques before the fusion stage.[15]

Genomic Encoder Architecture

The genomic subnetwork should be structured as an independent encoder that maps raw or pre-processed molecular profiles into a lower-dimensional latent embedding.[1] A hybrid approach utilizing 1D-Convolutional Neural Networks (1D-CNNs) is particularly effective for capturing local patterns in sequencing data, such as gene linkage or chromosomal proximity.[15] Alternatively, a pathway-aware Transformer-based encoder, such as the DeePathNet or SurvPath models, can be implemented.[19] This architecture groups individual genes into established biological pathways (e.g., cell cycle control, EMT, DNA repair) using priors from databases like MSigDB or KEGG.[20] By summarizing gene-level data into pathway-level embeddings, the model reduces the feature space while simultaneously enhancing interpretability.[20]

The proposed genomic encoder layers would include:

1. **Input Layer:** Standardized gene expression vectors (Log2 transformed and Z-score normalized).[23]
2. **Pathway Masking Layer:** A sparse linear layer that maps genes to their respective biological pathways, effectively performing a biologically-guided dimensionality reduction.[20]
3. **Transformer Blocks:** Multi-head self-attention mechanisms that learn the complex interactions between different biological pathways for a specific patient.[20]
4. **Embedding Layer:** A final dense layer that compresses the pathway features into a 128-dimension latent vector, matching the size of the image and clinical feature vectors to ensure balanced fusion.[1]

Variational Autoencoders (VAEs) for Latent Representation

Given the scarcity of sarcoma data, using a Variational Autoencoder (VAE) within the genomic subnetwork provides an additional layer of robustness.[15] A VAE can be pre-trained in an unsupervised manner on large pan-cancer datasets like the Cancer Genome Atlas (TCGA) to learn a generalizable representation of oncogenic genomic alterations.[26] The SarcomaAI tool can then utilize the encoder portion of this VAE as its genomic subnetwork, fine-tuning the weights on the sarcoma-specific cohort.[28] This transfer learning strategy leverages the billions of data points available in the broader oncology field to inform the analysis of rare mesenchymal tumors.[30]

Multimodal Fusion Paradigms and Cross-Modal Interaction

In the expanded SarcomaAI framework, the outputs from the three subnetworks—clinical, image, and genomic—must be integrated to produce the final prognostic evaluation.[1] While the existing SarcomaAI tool employs late fusion via concatenation, research indicates that intermediate fusion with cross-modal attention mechanisms can significantly enhance the model's ability to capture synergy between different modalities.[19]

Cross-Modal Attention and Interpretability

Cross-modal attention modules allow the network to "attend" to specific features in one modality based on signals from another.[19] For instance, a genomic signal indicating high chromosomal instability (via a high CINSARC score) could guide the image subnetwork to focus more intensely on areas of tumor necrosis or irregular vascularity in the MRI.[19] This interaction mimics the clinical workflow of a multidisciplinary tumor board, where radiologic findings are interpreted in the context of molecular pathology.[1]

The mathematical representation of this fusion in the model can be described by a cross-attention function:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

In this context, the query (Q) could be the genomic embedding, while the keys (K) and values (V) are derived from the feature maps of the image subnetwork.[19] This enables the model to create a "joint representation" that is more informative than the sum of its parts.[19]

Gating Networks and Mixture of Experts

To handle the extreme heterogeneity of sarcoma, the framework can implement a gating network or "mixture of multimodal experts" (MoME).[25] Not all modalities are equally informative for every patient. In a patient with a translocation-driven synovial sarcoma, the genomic subnetwork might carry the highest weight, whereas for a large, undifferentiated pleomorphic sarcoma, the image-derived volumetric and textural features might be more predictive.[25] A gating network dynamically calculates weights for each subnetwork's contribution to the final prediction based on the input data, ensuring that the model remains robust across the entire spectrum of STS subtypes.[25]

Fusion Strategy	Mathematical Mechanism	Clinical Rationale
Late Concatenation	$V_{fusion} = [V_{img} \oplus V_{clin} \oplus$	Simple integration of independent features.[1]
Cross-Attention	Genomic query attends to Image/Clinical keys	Molecular state context for morphological findings.[19]
Gating (MoME)	$\alpha_{img} V_{img} + \alpha_{clin} V_{clin} +$	Dynamic weighting based on tumor subtype info.[25]
Bilinear Pooling	$V_{fusion} = V_{img} \otimes$	Captures second-order multiplicative interactions.[17]

Scaling Gradient Blending for Three Modalities

The introduction of genomic data necessitates a more complex training objective to avoid modality competition and clash.[34] The gradient blending (GB) protocol must be scaled to account for the four-head output of the training phase: the global multimodal head and the three individual modality heads.[1]

The total loss function L_{total} is defined as:

$$L_{total}(t) = w_{MM}(t)L_{MM} + w_{img}(t)L_{img} + w_{clin}(t)L_{clin} + w_{gen}(t)L_{gen}$$

Where each $w_i(t)$ is the blending weight for modality i at epoch t, calculated based on the overfitting-to-generalization ratio.[1] The OGR is determined by comparing the training loss decrease to the validation loss increase over a set of epochs, identifying which branch is starting to overfit.[1] This is particularly vital for the genomic branch, which, due to its high dimensionality ($> 20,000$ features), is significantly more prone to overfitting than the image or clinical branches.[16] By penalizing the gradients of the genomic branch when its OGR rises, the model forces the shared layers to learn more generalizable representations from the molecular data.[1]

Furthermore, to combat "modality clash"—where the gradient update directions of the genomic branch conflict with those of the image branch—techniques like Gradient Orthogonalization and Adaptive Leveraging (GOAL) or Cross-Modality Gradient Harmonization should be integrated.[38] These methods ensure that the training steps for morphological feature extraction (MRI) and molecular pathway modeling (genetics) are mutually reinforcing rather than destructive.[38]

Data Preprocessing and Standardization Pipelines for Genomics

Just as MRI data requires N4 bias correction and Z-score normalization to ensure consistency across different scanners, genomic data requires a rigorous preprocessing pipeline to eliminate technical artifacts and batch effects.[1]

Transcriptomic Data Normalization

For RNA-seq data derived from the TCGA-SARC cohort or internal institutional databases, the pipeline must standardize the gene expression units.[23] The SarcomaAI playbook should be updated to include scripts for converting raw read counts into Transcripts Per Million (TPM) or Fragments Per Kilobase of transcript per Million mapped reads (FPKM), which account for both gene length and sequencing depth.[40] Following this, a logarithmic transformation (typically $Log2(x+1)$) is applied to reduce the influence of outliers and stabilize the variance.[23]

For multi-center data integration, the most critical step is batch effect correction.[23] Tools like ComBat or Feature Specific Quantile Normalization (FSQN) must be used to ensure that a patient's transcriptomic profile is comparable whether it was sequenced at Memorial Sloan Kettering (MSKCC) or McGill University.[23] This harmonization ensures that the neural network learns biological signals rather than institutional sequencing signatures.[23]

Copy Number Variation (CNV) Preprocessing

Processing CNV data for deep learning involves a transformation from discrete segments to a gene-level matrix.[41] The pipeline should filter for somatic mutations, discarding germline variations that do not contribute to the tumor's oncogenic progression.[41] Recurrent genomic alterations are identified using packages like GAIA, and aberrant regions are annotated with gene identifiers from Ensembl or NCBI.[41] The resulting data is typically a wide matrix where columns represent specific chromosomal gains or losses, which can then be flattened into a 1D vector for the genomic subnetwork.[40]

Preprocessing Step	Method/Protocol	Purpose
Quality Control	Detection of low-quality cells/reads	Eliminates noise from degraded samples.[16]
Unit Standardization	FPKM-UQ or TPM	Normalizes for sequencing depth and gene length.[40]
Variance Stabilization	Log2 Transformation	Reduces the weight of extreme outliers.[23]
Distribution Mapping	Quantile Normalization	Ensures identical distributions across cohorts.[23]
Feature Selection	Top-K Variable Genes	Reduces dimensionality for small-n datasets.[40]

Federated Learning for Multi-Institutional Genomic Collaboration

Building a robust multimodal model for rare diseases like sarcoma requires a scale of data that no single institution can provide.[1] However, the sharing of raw genomic sequences across borders is fraught with legal, ethical, and privacy challenges.[45] The SarcomaAI framework's use of federated learning (FL) via NVFlare provides a secure solution for integrating multi-institutional genomic data.[1]

Secure Aggregation and Privacy-Enhancing Technologies (PETs)

In the federated setting, the raw genomic data remains behind the institutional firewall of each participating site—whether at a university health center or a specialized cancer hospital.[1] Each institution hosts a client server that trains a local version of the SarcomaAI model.[1] At the end of each training epoch, only the updated weights and gradients of the genomic subnetwork are transmitted to the central server managed by the project leads.[48]

To provide maximum privacy guarantees, the framework incorporates several PETs:

1. **Secure Aggregation:** The central server receives encrypted model updates from multiple clients and aggregates them (e.g., via FedAvg) before decryption, ensuring that

no single site's genomic gradients are visible in isolation.[29]

2. **Differential Privacy (DP):** A controlled amount of Gaussian noise is added to the gradients before transmission. This ensures that the model cannot "memorize" specific rare mutations from a single patient, which could otherwise be reverse-engineered through a model inversion attack.[49]
3. **Homomorphic Encryption (HE):** Advanced cryptographic protocols that allow mathematical operations to be performed directly on encrypted data, providing a theoretical guarantee that sensitive molecular sequences are never exposed during transit or processing.[49]

Implementation on Cloud Infrastructure

The SarcomaAI playbook outlines a scalable AWS-based architecture for deploying this federated genomic pipeline.[1] Each institution utilizes a Virtual Private Cloud (VPC) with a private subnet to host their NVFlare client.[1] Genomic data, pre-processed into AI-ready formats like H5AD or wide CSVs, is stored in a private S3 bucket with strict IAM roles that only allow the local ML workstation access.[1] Communication with the central orchestrator is facilitated via the gRPC protocol through an AWS Transit Gateway, providing a high-performance, secure bridge for the model updates.[1]

Integrating Biological Priors and Clinical Signatures

Rather than relying entirely on the neural network to discover patterns in raw genomic data, the SarcomaAI tool can be significantly enhanced by integrating established clinical signatures as biological "priors".[21] This hybrid approach combines the strengths of traditional biomarker research with the pattern-recognition capabilities of deep learning.[14]

The CINSARC Signature as a Model Input

The 67-gene CINSARC signature is one of the most powerful prognostic markers in sarcoma research.[11] Developed by analyzing sarcomas with complex genetics, it targets genes involved in chromosome integrity and mitotic control (e.g., $AURKA$).[11] Patients can be stratified into C1 (low risk) and C2 (high risk) groups based on their CINSARC expression levels, which correlate strongly with metastatic relapse-free survival.[55] By including the CINSARC score as a high-weight feature in the genomic subnetwork, the SarcomaAI model can leverage a decade of validated biological insight.[11] This acts as a regularizer, ensuring the model's predictions align with known mechanisms of tumor aggressiveness.[21]

The Genomic Grade Index (GGI)

The Genomic Grade Index (GGI) is another proliferation-based signature, originally developed for breast cancer, that has been successfully adapted for sarcoma prognosis.[58] The GGI can refine the prediction of metastasis-free survival in patients who have undergone surgery, providing a molecular counterpoint to the histological FNCLCC grade.[58] Integrating GGI as a continuous variable alongside the discrete FNCLCC grade in the clinical subnetwork allows the model to identify cases where morphological grade is misleading—such as "grade 2" tumors that behave with the biological aggression of "grade 3" disease.[1]

Signature Name	Number of Genes	Key Biological Pathway	Clinical Application
CINSARC	67	Mitosis and Chromosome Integrity	Predicting metastatic relapse in complex-karyotype STS.[11]
GGI	97	Cell Proliferation	Refinement of pathological grading for survival.[58]
ICR	20	Cytotoxic Immune Response	Predicting response to immunotherapy and survival.[58]
SPM6	6 (Proteomic)	DNA Replication	Added to Sarculator for risk stratification.[60]

Interpretability and the Future of Subtype-Specific SarcomaAI

For a multimodal neural network to be adopted into clinical practice, it must transcend the "black box" nature of deep learning.[14] Clinicians require a clear understanding of why a model assigned a specific risk score to a patient. The integration of genomics provides new avenues for this transparency.[5]

Multimodal Saliency Visualization

The current SarcomaAI tool uses Grad-CAM to visualize areas of the tumor volume that contribute most to the prognosis prediction.[1] With the genomic subnetwork, this can be expanded into a unified multimodal salience map.[19] Using attention visualization or SHAP-based feature importance, the model can generate a report for the multidisciplinary tumor board that highlights:

- **Radiological Salience:** Areas of MRI heterogeneity or edema indicating high aggression.[1]
- **Genomic Salience:** Biological pathways (e.g., "Upregulation of Hypoxia markers" or "Downregulation of Apoptotic signaling") that drive the metastatic risk.[14]
- **Clinical Salience:** The specific impact of factors like patient age or tumor volume on the final score.[14]

Evolution toward Predictive Modeling for Targeted Therapy

While the current objective of SarcomaAI is to predict survival and metastasis, the inclusion of genomic data enables the model to evolve into a predictive tool for treatment response—a cornerstone of precision oncology.[1] Certain sarcoma subtypes harbor actionable mutations that are invisible to MRI.[59] For example, mutations in the IDH (isocitrate dehydrogenase) gene in chondrosarcomas or KIT and $PDGFRA$ mutations in gastrointestinal stromal tumors (GIST) define susceptibility to specific targeted inhibitors.[12] By training the genomic subnetwork to identify these molecular landscapes, SarcomaAI can help clinicians move beyond generic doxorubicin-based regimens and toward therapy tailored to the patient's unique genomic profile.[22]

Nuanced Challenges in Genomic Integration

As the SarcomaAI project implements these architectural changes, several second-order challenges must be considered.

Handling Missing Modalities

In real-world clinical settings, not all patients will have complete genomic profiling.[19] A truly robust multimodal system must be able to generate accurate predictions even when one modality is missing.[26] Modular architectures like MultiSurv handle this by using dedicated "missingness" masks and dropout-aware fusion layers.[21] If genomic data is unavailable, the model can fall back to its image and clinical branches, perhaps utilizing a pre-trained "imputation" branch that attempts to estimate key genomic signatures (like CINSARC) from the MRI radiomics features—a concept known as "radiogenomics".[12]

Small-n vs. Large-d

The "small-n, large-d" problem—where the number of genomic features (d) is significantly larger than the number of patients (n)—is a pervasive issue in rare disease research.[13] To overcome this, the SarcomaAI project must prioritize feature selection methods that identify the most variable and informative genes for sarcoma specifically.[40] Agglomerative hierarchical clustering and deterministic sampling based on leverage scores can be used to reduce the genomic input from 20,000 genes to a core set of 100-500 highly informative markers before feeding them into the neural network.[44]

Roadmap for Implementation: Version 0.2 of the SarcomaAI Pipeline

To achieve the integration of genomic data, the following strategic steps are proposed for the next version of the SarcomaAI playbook.[1]

Phase 1: Retrospective Multi-Omic Mapping

Institutions participating in the SarcomaAI project should begin by linking their existing imaging cohorts with available genomic data from internal pathology departments or public repositories like the GDC Data Portal.[1] This requires a robust mapping of patient identifiers to connect 3D MRI volumes with the corresponding genomic VCF or FASTQ files while maintaining absolute anonymity.[1]

Phase 2: Pre-training the Multi-Platform Encoder

The genomic subnetwork should be pre-trained using the Pan-Cancer Atlas or similar large-scale datasets.[20] This pre-training allows the network to learn the "language of genomics" before it is tasked with identifying the subtle signatures of rare sarcomas.[28] Simultaneously, the image subnetwork can be further refined using self-supervised learning on large orthopedic imaging datasets to improve its morphological feature extraction.[1]

Phase 3: Federated Deployment and Subtype-Specific Refinement

The final phase involves the deployment of the integrated clinical-image-genomic MMNN across the federated network using NVFlare.[1] As the pool of patients grows, the model can move away from "lumping" all sarcomas together and begin providing subtype-specific survival curves.[1] This will eventually allow for the development of "digital twins" in sarcoma care, where the AI provides a simulated outcome for various treatment scenarios based on the patient's integrated multimodal profile.[29]

Nuanced Conclusions on Architectural Evolution

The evolution of SarcomaAI from a clinical-radiological tool into a comprehensive multimodal clinical-genomic framework represents a transformative milestone in orthopaedic oncology. The inclusion of genomic data—specifically copy number variations and transcriptomic signatures like CINSARC—addresses the fundamental biological heterogeneity that has long limited the accuracy of traditional prognostic models.[1] By employing a pathway-aware genomic subnetwork and advanced gradient blending protocols, the architecture can handle the disparate dimensionalities of its inputs while avoiding the pitfalls of modality competition and clash.[9]

The success of this integration relies heavily on the federated learning infrastructure already established in the SarcomaAI playbook. By leveraging NVFlare and trusted execution environments, the project can aggregate the critical mass of genomic data required for rare disease modeling while ensuring that sensitive patient information never leaves its institution of origin.[1] As the field moves toward a future defined by precision medicine, the synergy of morphological imaging and molecular biology, facilitated by the expanded SarcomaAI framework, provides a clear roadmap for improving the lives of patients with these complex and challenging tumors. The transition from aggregate median statistics to individualized, biologically-grounded risk assessment marks the true arrival of precision oncology in the treatment of soft tissue sarcomas.

Works cited

1. SarcomaAI-Playbook_v0.1.pdf
2. Transcriptomic Profiling of Old Age Sarcoma Patients using TCGA RNA-seq data - bioRxiv, accessed March 11, 2026, https://www.biorxiv.org/content/10.1101/2025.01.03.631189v1.full.pdf
3. Sarcoma Study - NCI - National Cancer Institute, accessed March 11, 2026, https://www.cancer.gov/ccg/research/genome-sequencing/tcga/studied-cancers/sarcoma-study
4. Comprehensive and Integrated Genomic Characterization of Adult ..., accessed March 11, 2026, https://gdc.cancer.gov/about-data/publications/sarc_2017
5. AI-Driven Early Detection of Sarcoma: A Machine Learning and Deep Learning-Based Approach., accessed March 11, 2026, https://spast.org/techrep/article/download/5498/638/10833
6. Balanced Multimodal Learning via On-the-fly Gradient Modulation - People, accessed March 11, 2026, https://people.cs.vt.edu/chris/cs6804_spring2023/slides/connor_learning.pdf
7. Optimizing Multi-task Learning Models in Practice | Towards Data Science, accessed March 11, 2026, https://towardsdatascience.com/optimizing-multi-task-learning-models-in-practice-bde4f18f0bd8/
8. What is Gradient Blending and focal loss | by Anandut | Medium, accessed March 11, 2026, https://medium.com/@anandut2001/what-is-gradient-blending-and-focal-loss-2ac2bbc90801
9. Towards Optimal Multi-Modal Federated Learning on Non-IID Data with Hierarchical Gradient Blending - iQua, accessed March 11, 2026, https://iqua.ece.toronto.edu/papers/schen-infocom22.pdf
10. Clinical Data - cBioPortal for Cancer Genomics, accessed March 11, 2026,

http://www.cbioportal.org/study/clinicalData?id=sarc_tcga_pub
11. A Global and Integrated Analysis of CINSARC-Associated Genetic Defects - AACR Journals, accessed March 11, 2026, https://aacrjournals.org/cancerres/article/80/23/5282/645883/A-Global-and-Integrated-Analysis-of-CINSARC
12. Prognostic prediction in soft-tissue sarcomas using deep learning and digital pathology of tumor and margin areas - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12586537/
13. Evaluating dimensionality reduction for genomic prediction - Frontiers, accessed March 11, 2026, https://www.frontiersin.org/journals/genetics/articles/10.3389/fgene.2022.958780/full
14. Multimodal feature-optimized approaches for cancer classification using microarray gene expression analysis - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12612085/
15. A deep learning framework for characterization of genotype data - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC8896001/
16. Single-cell RNA sequencing data dimensionality reduction (Review) - Spandidos Publications, accessed March 11, 2026, https://www.spandidos-publications.com/10.3892/wasj.2025.315
17. Deep learning approaches for resolving genomic discrepancies in cancer: a systematic review and clinical perspective | Briefings in Bioinformatics | Oxford Academic, accessed March 11, 2026, https://academic.oup.com/bib/article/26/6/bbaf541/8312888
18. Dimensionality Reduction of Genetic Data using Contrastive Learning - bioRxiv.org, accessed March 11, 2026, https://www.biorxiv.org/content/10.1101/2024.09.30.615901v1.full
19. Multimodal deep learning for cancer prognosis prediction with clinical information prompts integration - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12830675/
20. Transformer-based deep learning integrates multi-omic data with cancer pathways | bioRxiv, accessed March 11, 2026, https://www.biorxiv.org/content/10.1101/2022.10.27.514141v2.full-text
21. Structural Prognostic Event Modeling for Multimodal Cancer Survival Analysis - arXiv.org, accessed March 11, 2026, https://arxiv.org/html/2512.01116v2
22. An interpretable deep learning framework for genome-informed precision oncology - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC10369905/
23. A comparison of RNA-Seq data preprocessing pipelines for transcriptomic predictions across independent studies - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC11080237/
24. Self-Normalizing Deep Learning for Enhanced Multi-Omics Data Analysis in Oncology, accessed March 11, 2026, https://www.preprints.org/manuscript/202504.0534
25. MoME: Mixture of Multimodal Experts for Cancer Survival Prediction | MICCAI

2024, accessed March 11, 2026, https://papers.miccai.org/miccai-2024/531-Paper2168.html
26. Long-term cancer survival prediction using multimodal deep learning - PMC - NIH, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC8242026/
27. Multimodal Dimension Reduction and Subtype Classification of Head and Neck Squamous Cell Tumors - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC9326399/
28. Federated learning-based protein language models with Apheris on AWS, accessed March 11, 2026, https://aws.amazon.com/blogs/industries/federated-learning-based-protein-language-models-with-apheris-on-aws/
29. OmniFed: A Modular Framework for Configurable Federated Learning from Edge to HPC, accessed March 11, 2026, https://arxiv.org/html/2509.19396v1
30. A machine learning approach for multimodal data fusion for survival prediction in cancer patients - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12053085/
31. MULTIMODAL MACHINE LEARNING TO ENHANCE PATIENT PROGNOSIS & MANAGEMENT IN CANCER CARE - Dartmouth Digital Commons, accessed March 11, 2026, https://digitalcommons.dartmouth.edu/cgi/viewcontent.cgi?article=1424&context=dissertations
32. Page-Net: Interpretable and Integrative Deep Learning for Survival Analysis Using Histopathological Images and Genomic Data - OAsis Research Repository - University of Nevada, Las Vegas, accessed March 11, 2026, https://oasis.library.unlv.edu/compsci_fac_articles/223/
33. Structured Multimodal Deep Learning improves Genomic Prediction in Future Environments, accessed March 11, 2026, https://www.biorxiv.org/content/10.1101/2025.09.05.674546v1.full-text
34. Boosting Multimodal Learning via Disentangled Gradient Learning - arXiv, accessed March 11, 2026, https://arxiv.org/html/2507.10213v1
35. Improving Multimodal Learning Balance and Sufficiency through Data Remixing - arXiv, accessed March 11, 2026, https://arxiv.org/html/2506.11550v2
36. Improving Multimodal Learning Balance and Sufficiency through Data Remixing - GitHub, accessed March 11, 2026, https://raw.githubusercontent.com/mlresearch/v267/main/assets/ma25c/ma25c.pdf
37. A novel modality contribution confidence-enhanced multimodal deep learning framework for multiomics data - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12577108/
38. Scaling Multimodal Pre-Training via Cross-Modality Gradient Harmonization - NeurIPS, accessed March 11, 2026, https://proceedings.neurips.cc/paper_files/paper/2022/hash/eacad5b8e67850f2b8dd33d87691d097-Abstract-Conference.html
39. GOAL: Balance Multimodal Learning with Gradient Orthogonalization and Adaptive Leveraging | OpenReview, accessed March 11, 2026,

https://openreview.net/forum?id=I3uFqoUZ2Y
40. TCGA Genomic Data Commons ETL pipelines - Matillion Exchange, accessed March 11, 2026, https://exchange.matillion.com/data-productivity-cloud/pipeline/tcga-genomic-data-commons/
41. MLOmics: Benchmark for Machine Learning on Cancer Multi-Omics Data - arXiv, accessed March 11, 2026, https://arxiv.org/html/2409.02143v2
42. Mastering Data Preprocessing for Neural Networks | CodeSignal Learn, accessed March 11, 2026, https://codesignal.com/learn/courses/introduction-to-neural-networks-with-tensorflow/lessons/mastering-data-preprocessing-for-neural-networks
43. emdann/scATAC_prep: A not-too-orderly collection of scripts for preprocessing and dimensionality reduction of scATAC-seq data - GitHub, accessed March 11, 2026, https://github.com/emdann/scATAC_prep
44. Evaluating dimensionality reduction for genomic prediction - PMC - NIH, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC9614092/
45. NVIDIA FLARE Day 2025, accessed March 11, 2026, https://developer.nvidia.com/flare-day-2025
46. What is Federated Learning? — NVIDIA FLARE 2.4.0 documentation, accessed March 11, 2026, https://nvflare.readthedocs.io/en/2.4/fl_introduction.html
47. A Pragmatic Framework for Federated Learning Risk and Governance in Academic Medical Centers - JMIR AI, accessed March 11, 2026, https://ai.jmir.org/2026/1/e80022
48. NVIDIA FLARE | NVIDIA Developer, accessed March 11, 2026, https://developer.nvidia.com/flare
49. NVIDIA FLARE Overview - Read the Docs, accessed March 11, 2026, https://nvflare.readthedocs.io/en/main/flare_overview.html
50. On the Security and Privacy of Federated Learning: A Survey with Attacks, Defenses, Frameworks, Applications, and Future Directions - arXiv.org, accessed March 11, 2026, https://arxiv.org/html/2508.13730v1
51. [2210.13291] NVIDIA FLARE: Federated Learning from Simulation to Real-World - ar5iv, accessed March 11, 2026, https://ar5iv.labs.arxiv.org/html/2210.13291
52. Security in NVIDIA FLARE Federated Computing Systems - GitHub, accessed March 11, 2026, https://github.com/NVIDIA/NVFlare/blob/main/examples/tutorials/self-paced-training/part-3_security_and_privacy/chapter-6_Security_in_federated_compute_system/06.0_introduction/introduction.ipynb
53. A Pragmatic Framework for Federated Learning Risk and Governance in Academic Medical Centers - JMIR AI, accessed March 11, 2026, https://ai.jmir.org/2026/1/e80022/PDF
54. (PDF) The CINSARC signature as a prognostic marker for clinical outcome in multiple neoplasms - ResearchGate, accessed March 11, 2026, https://www.researchgate.net/publication/318460631_The_CINSARC_signature_as_a_prognostic_marker_for_clinical_outcome_in_multiple_neoplasms
55. CINSARC and Sarculator in Patients with Primary Retroperitoneal Sarcoma: A

Combined Analysis of Single-Institution Data and the EORTC-STBSG-62092 Trial (STRASS) - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12314516/
56. The CINSARC signature predicts the clinical outcome in patients with Luminal B breast cancer - PMC, accessed March 11, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC8099860/
57. Complexity index in sarcoma and genomic grade... : Pediatric Blood & Cancer - Ovid, accessed March 11, 2026, https://www.ovid.com/journals/pedbc/fulltext/10.1002/pbc.28987~complexity-index-in-sarcoma-and-genomic-grade-index-gene
58. Immunologic constant of rejection signature is prognostic in soft-tissue sarcoma and refines the CINSARC signature - The Journal for ImmunoTherapy of Cancer, accessed March 11, 2026, https://jitc.bmj.com/content/10/1/e003687
59. Transcriptomic-Based Classification Identifies Prognostic Subtypes and Therapeutic Strategies in Soft Tissue Sarcomas - MDPI, accessed March 11, 2026, https://www.mdpi.com/2072-6694/17/17/2861
60. Professor Paul Huang - The Institute of Cancer Research, accessed March 11, 2026, https://www.icr.ac.uk/research-and-discoveries/find-a-researcher/test-researcher-profile-detail/dr-paul-huang
61. High-Dimensional Cancer Genomic Data Modeling Using Distributed Machine Learning Algorithms for Genetic and Transcriptomic Pattern Discovery, accessed March 11, 2026, https://geneticsmr.com/index.php/gmr/article/view/314
62. Gene Expression Data-Based Interpretable Machine Learning Framework for Classifying Brain Cancer Subtypes - Biomedical and Pharmacology Journal, accessed March 11, 2026, https://biomedpharmajournal.org/vol18no3/gene-expression-data-based-interpretable-machine-learning-framework-for-classifying-brain-cancer-subtypes/
63. Machine learning-based analysis of genomic and transcriptomic data unveils sarcoma clusters with superlative prognostic and predictive value - medRxiv.org, accessed March 11, 2026, https://www.medrxiv.org/content/10.1101/2025.01.31.25321492v1.full-text
64. Professor Robin Jones - The Institute of Cancer Research, accessed March 11, 2026, https://www.icr.ac.uk/research-and-discoveries/find-a-researcher/test-researcher-profile-detail/dr-robin-jones
65. Comprehensive Molecular Characterization of Soft Tissue Sarcoma for Prediction of Pazopanib-Based Treatment Response - Cancer Research and Treatment, accessed March 11, 2026, https://e-crt.org/journal/view.php?number=3407
66. TCGA-SARC - The Cancer Imaging Archive (TCIA), accessed March 11, 2026, https://www.cancerimagingarchive.net/collection/tcga-sarc/
67. TCGADownloadHelper: simplifying TCGA data extraction and preprocessing - Frontiers, accessed March 11, 2026, https://www.frontiersin.org/journals/genetics/articles/10.3389/fgene.2025.1569290/full

68. Secure Federated Learning with NVIDIA & Google Cloud - Duality Tech, accessed March 11, 2026, https://dualitytech.com/blog/secure-federated-learning-with-nvidia-and-google-cloud/